

Karlsruhe Machine Learning,
Statistics and AI Meetup

Introduction to TensorFlow with Application to Autonomous Driving



Christian Hubschneider
FZI Forschungszentrum Informatik

THIS IS YOUR MACHINE LEARNING SYSTEM?

| YUP! YOU POUR THE DATA INTO THIS BIG
PILE OF LINEAR ALGEBRA, THEN COLLECT
THE ANSWERS ON THE OTHER SIDE.

| WHAT IF THE ANSWERS ARE WRONG? |

JUST STIR THE PILE UNTIL
THEY START LOOKING RIGHT.



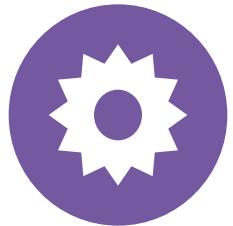
Topics



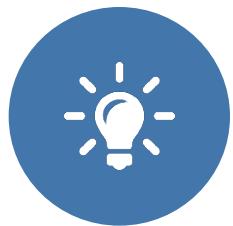
Overview
over
Frameworks



TensorFlow
Basics



Coding
Examples



TensorFlow
for Autonomous
Driving



How to get
involved



Deep Learning Frameworks

OVERVIEW OVER CROSS-PLATFORM MACHINE LEARNING TOOLS

- Caffe: caffe.berkeleyvision.org/ - U.C. Berkeley
 - C++, Python, MATLAB
- Caffe2: caffe2.ai/ - Facebook
 - C++, Python
 - Released: April 18th 2017: <https://caffe2.ai/blog/2017/04/18/caffe2-open-source-announcement.html>
- MXNet: mxnet.io/ - Amazon / AWS
 - C++, Python, Julia, Scala, R, MATLAB, ..
- TensorFlow: <https://www.tensorflow.org/> - Google
 - C/C++, Python, Java, Go
- Theano: <http://wwwdeeplearning.net/software/theano/>
 - Python
- Torch: <http://torch.ch/> - New York University
 - C/C++, Lua, LuaJIT
- Keras, Lasagne, ...

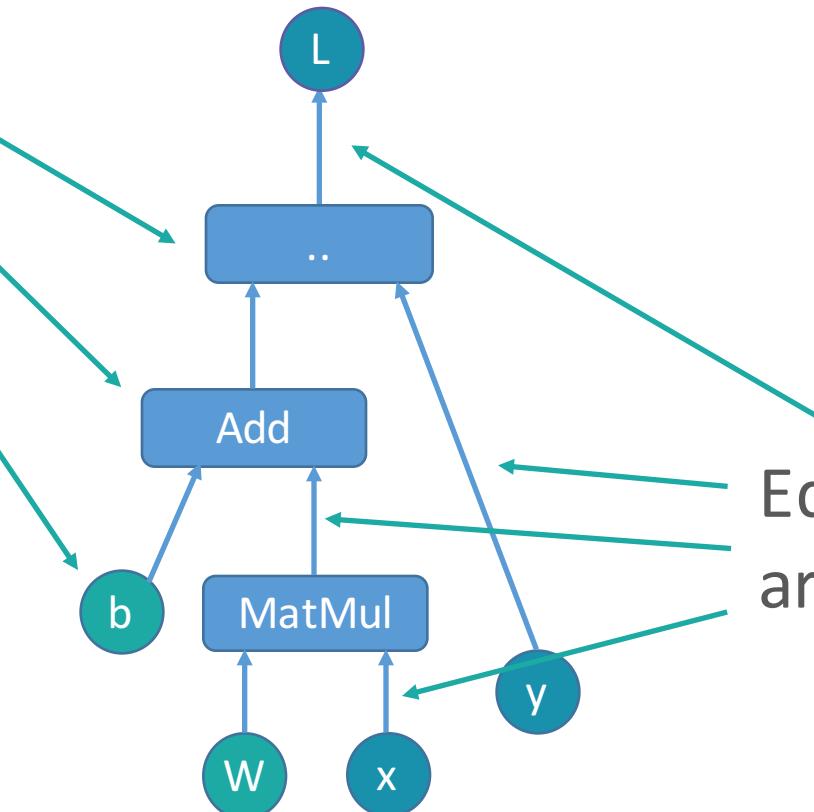




A Simple Graph

TENSORFLOW REPRESENTS THE LEARNING PROBLEM AS A GRAPH

Nodes are **Operations**
Each node has zero or
more inputs and zero or
more outputs.



Edges are n-dimensional
arrays: **Tensors**

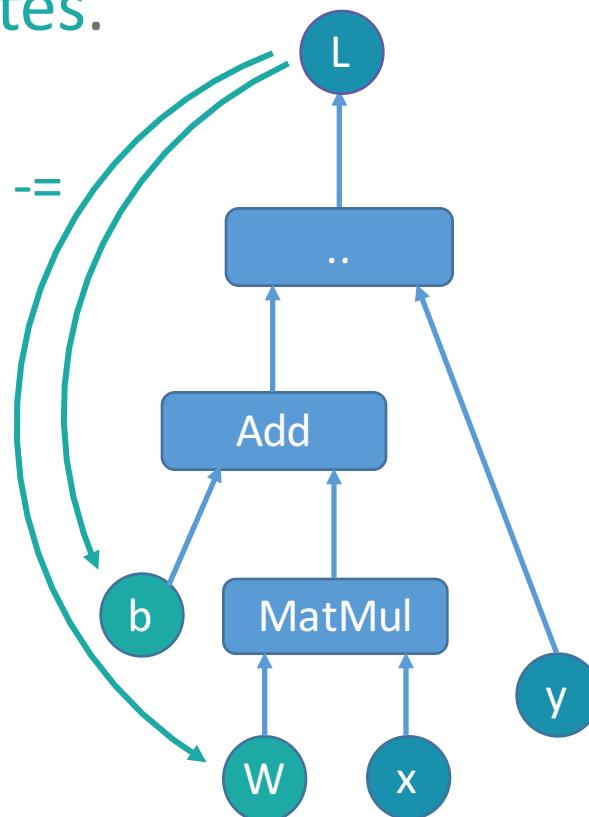


A Simple Graph

NUMERICAL COMPUTATION USING DATA FLOW GRAPHS

The graph has a **state** that can be used to perform gradient descent (or any other) **weight updates**.

Update bias
and weights.

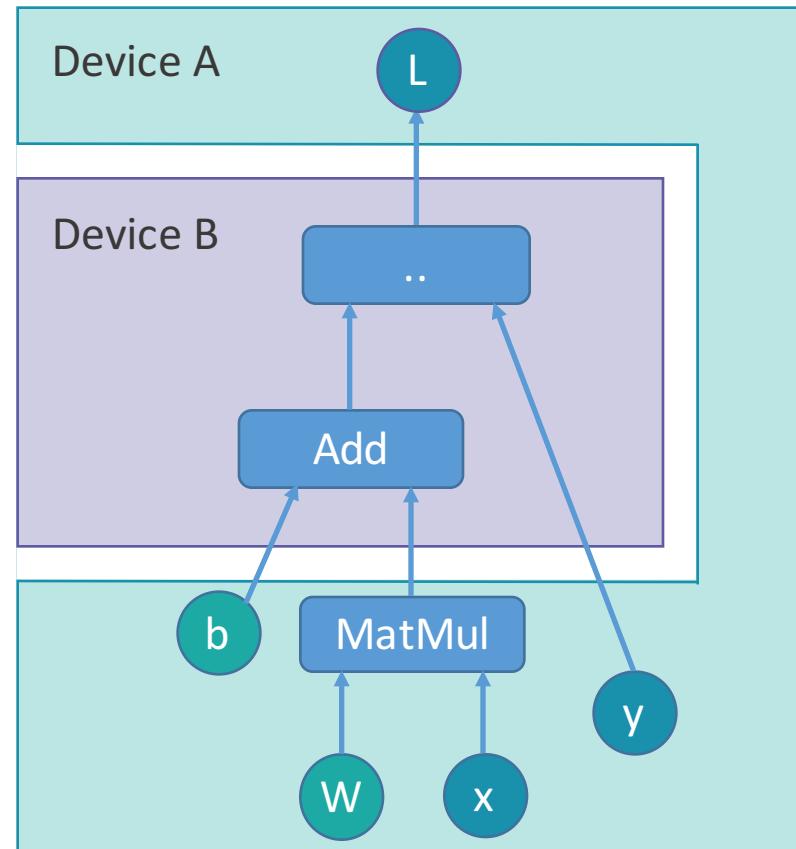




A Simple Graph

SCALABILITY THROUGH THE GRAPH STRUCTURE

In TensorFlow the graph makes **distributed processing** on multiple devices easy.



Devices: Processes,
Machines, CPUs, GPUs, ..



Overview over Operators

OPERATIONS REPRESENT ABSTRACT COMPUTATIONS ON A NUMBER OF NODES

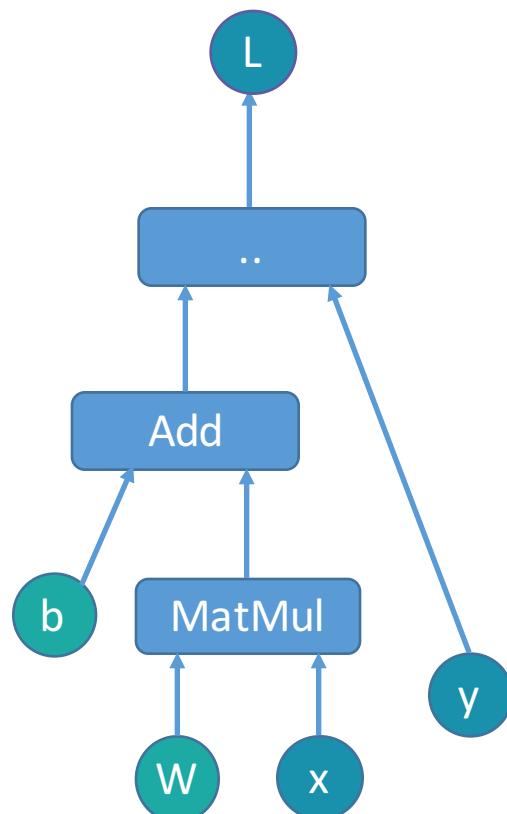
Category	Example
Element-wise mathematical operations	Add, Sub, Mul, Div, Exp, Log, Compare, ..
Array operations	Concat, Slice, Split, Shuffle, ..
Matrix operations	MatMul, MatrixInverse, MatrixDeterminant, ..
Stateful operations	Variable, Assign, AssignAdd, ..
Neural-net building blocks (layers, non-linearities)	SoftMax, Sigmoid, ReLU, Conv2D, MaxPool, ..
Checkpoint operations	Save, Restore
Queue and sync operations	Enqueue, Dequeue, ..
Control flow operations	Merge, Switch, Enter, Leave, NextIteration



A Simple Graph

NUMERICAL COMPUTATION USING DATA FLOW GRAPHS

A graph is defined in Python by attaching nodes to the **global graph**.



```
import tensorflow as tf

# Placeholders for input
x = tf.placeholder(tf.float32)
y = tf.placeholder(tf.float32)

# Variables to be learned
w = tf.get_variable("W", [1, 1])
b = tf.get_variable("b", [1])

y_pred = tf.matmul(x, w) + b
loss = tf.reduce_mean(tf.square(y - y_pred))
```

Nothing is
executed yet!

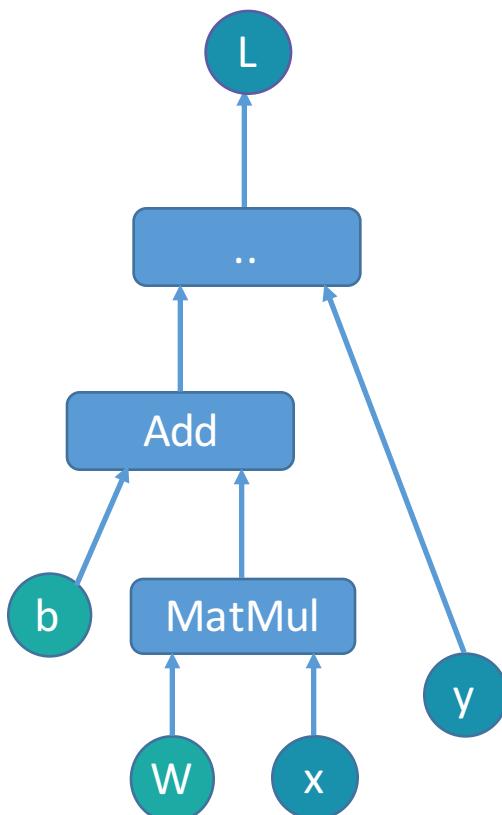
$$L(W, b) = \frac{1}{N} \sum_{i=1}^N (y_i - (Wx_i + b))^2$$



Tensorflow Sessions

NUMERICAL COMPUTATION USING DATA FLOW GRAPHS

Clients interact with and run the TensorFlow engine using **Sessions**.

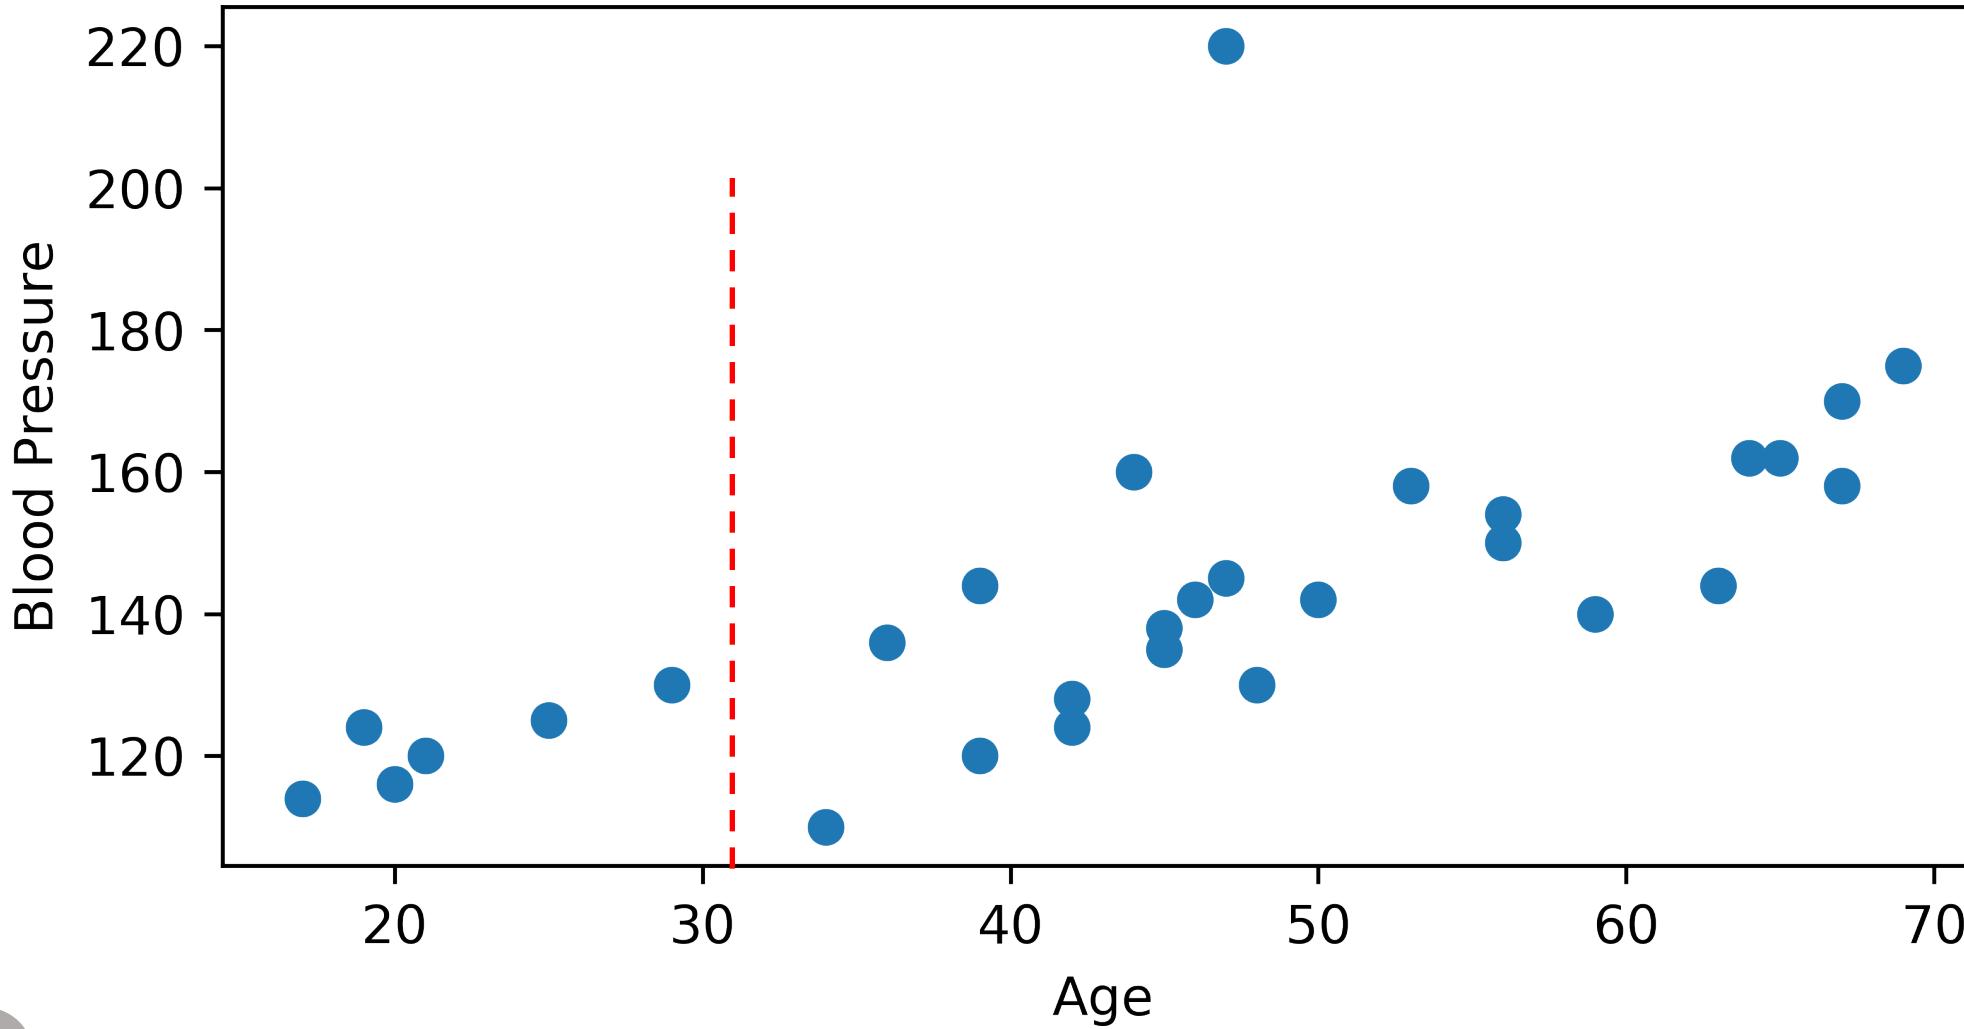


```
training_step =  
    tf.train.GradientDescentOptimizer(0.001).minimize(loss)  
  
with tf.Session() as sess:  
    # Initialize all variables in graph  
    sess.run(tf.global_variable_initializer())  
  
    xs, ys = data[0], data[1] # choose data for step  
  
    for _ in range(10000):  
        # Run gradient descent step  
        sess.run(training_step, feed_dict={x: xs, y: ys})
```



Linear Regression

WHAT SHOULD MY BLOOD PRESSURE BE WHEN I AM 31?





Linear Regression

WHAT SHOULD MY BLOOD PRESSURE BE WHEN I AM 31?

Assume a model

$$y = aX + b + \epsilon,$$

where a and b are unknown constants that represent **slope** and **intercept** and error term ϵ .

Given some estimates \hat{a} and \hat{b} for the model coefficients, we can predict the blood pressure for arbitrary ages X .

How can we find those estimates? (e.g. using Gradient Descent*)

*Be aware that for this simple case way better solutions than Gradient Descent do exist! ;-)



Simple Linear Regression using TensorFlow

WHAT SHOULD MY BLOOD PRESSURE BE WHEN I AM 31?

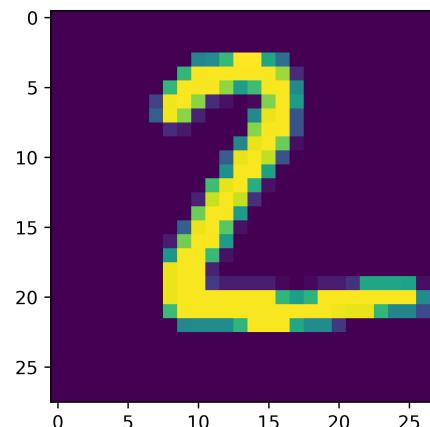
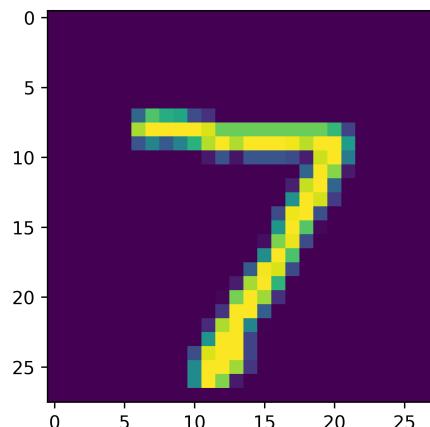
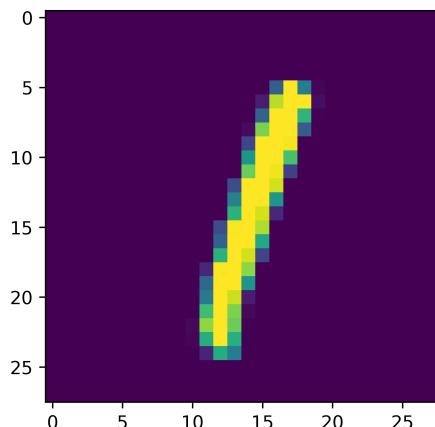
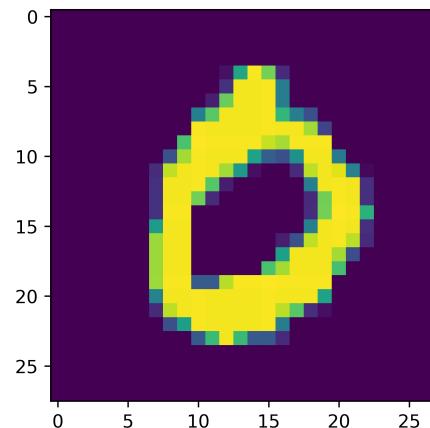
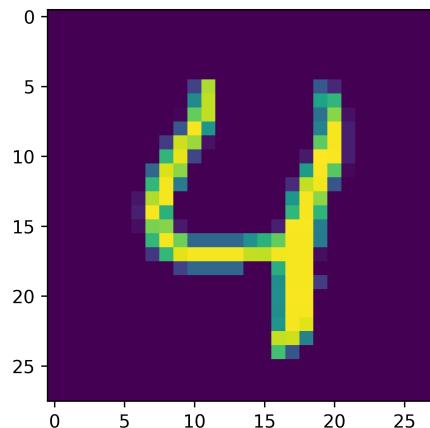
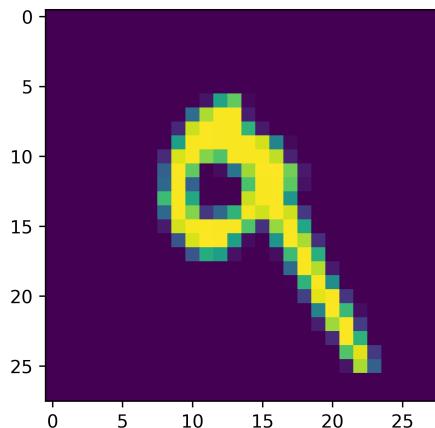
Live demo

`linear_regression.ipynb`



The MNIST Dataset

THE MNIST DATABASE OF HANDWRITTEN DIGITS (by Yann LeCun et al.)



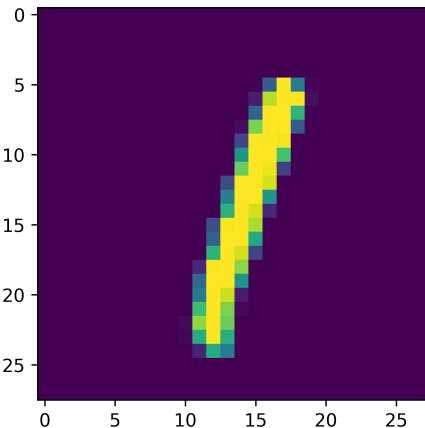
Dataset containing handwritten digits (0-9)

- 60'000 training samples
- 10'000 test samples
- Resolution: 28x28
- State of the art:
0.21% classification error rate



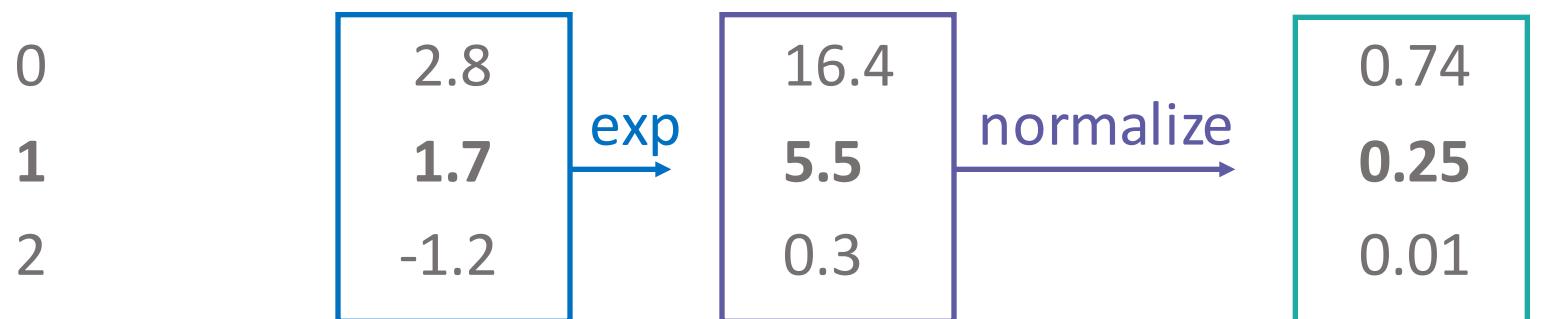
Logistic Regression, Softmax and Classification

CLASSIFICATION INSTEAD OF REGRESSION



$$L_i = -\log \frac{e^{s_{y_i}}}{\sum_j e^{s_j}}$$

unnormalized probabilities



unnormalized log
probabilities

probabilities

Softmax Loss:
 $L_i = -\log(0.25) = 0.61$



„Multinomial Logistic Regression“ on MNIST

USING LOGISTIC REGRESSION TO SOLVE A CLASSIFICATION TASK ON MNIST

Live demo

[softmax_classification_mnist.ipynb](#)

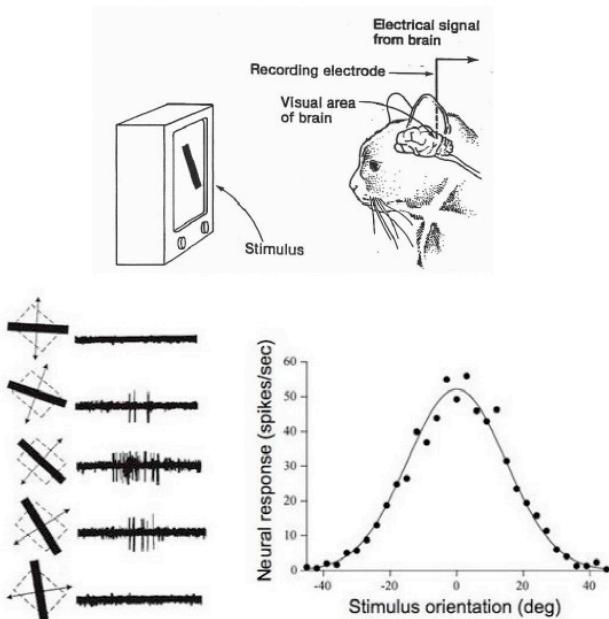


Convolutional Neural Networks

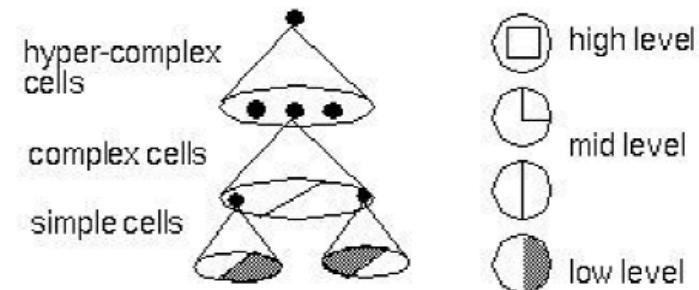
OR HOW TO GET EVEN BETTER CLASSIFICATION RESULTS

How to get even better classification results?

Use **spatial relationships** between pixels in the image (inspired by Biology)!



featural hierarchy



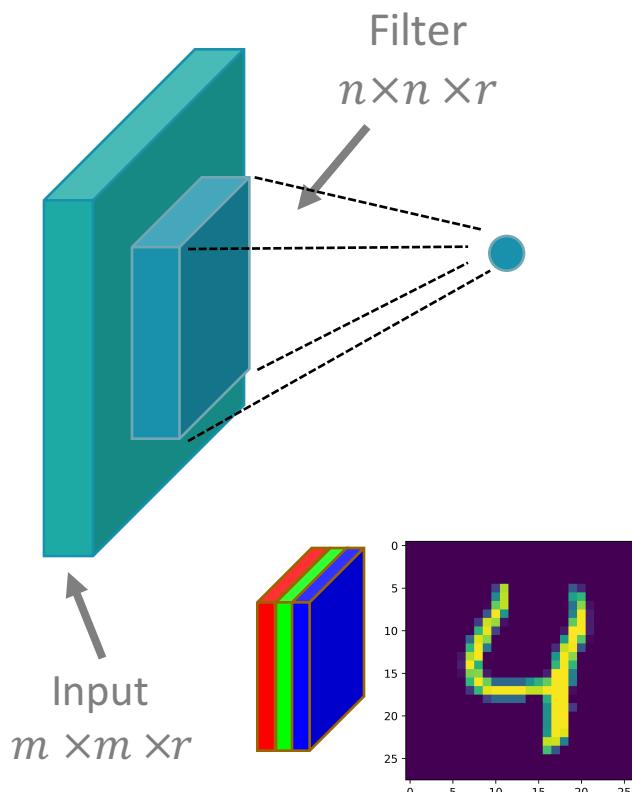
[Hubel & Wiesel 1959, 1968]



Building Blocks of a CNN: Convolutional Layer

OR HOW TO GET EVEN BETTER CLASSIFICATION RESULTS

Each neuron only sees part of the input (“receptive field”).

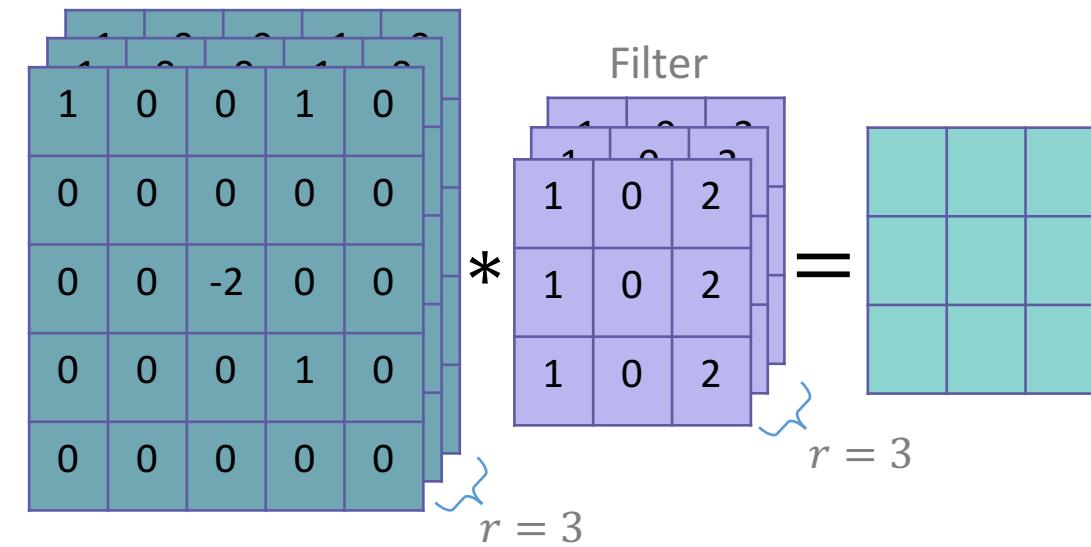
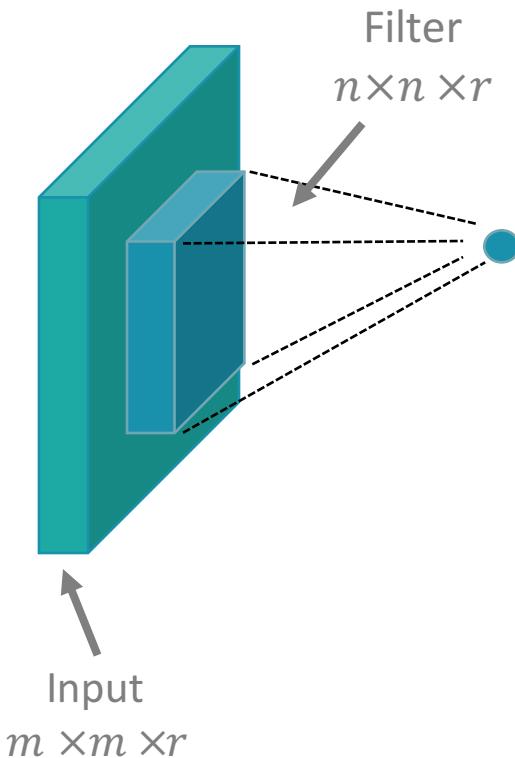




Building Blocks of a CNN: Convolutional Layer

OR HOW TO GET EVEN BETTER CLASSIFICATION RESULTS

Output of the neuron: **discrete convolution** of all input-features for one filter.



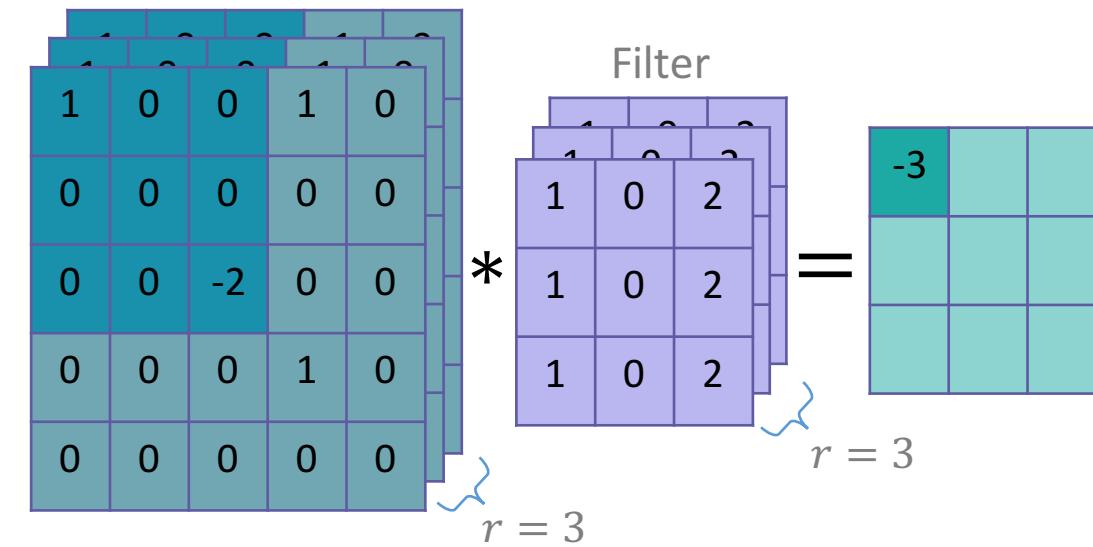
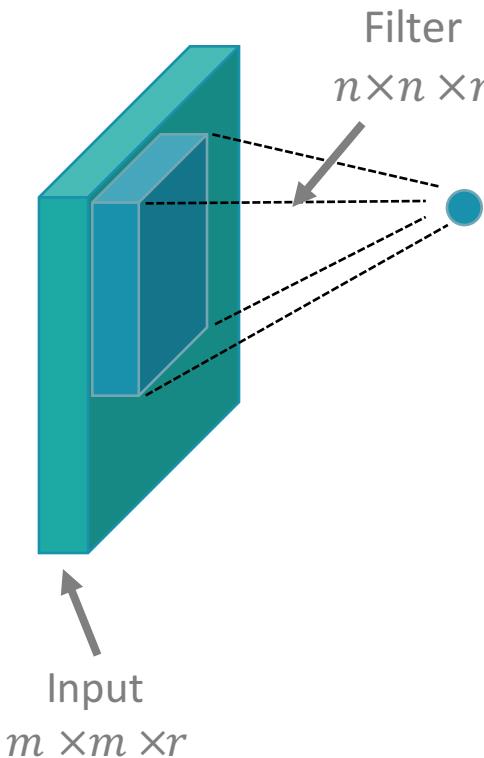
$$(I * F)_{r,s} := \sum_{u=0}^{r-1} \sum_{v=0}^{r-1} I_{r+u,s+v} \cdot F_{u,v}$$



Building Blocks of a CNN: Convolutional Layer

OR HOW TO GET EVEN BETTER CLASSIFICATION RESULTS

Output of the neuron: **discrete convolution** of all input-features for one filter.



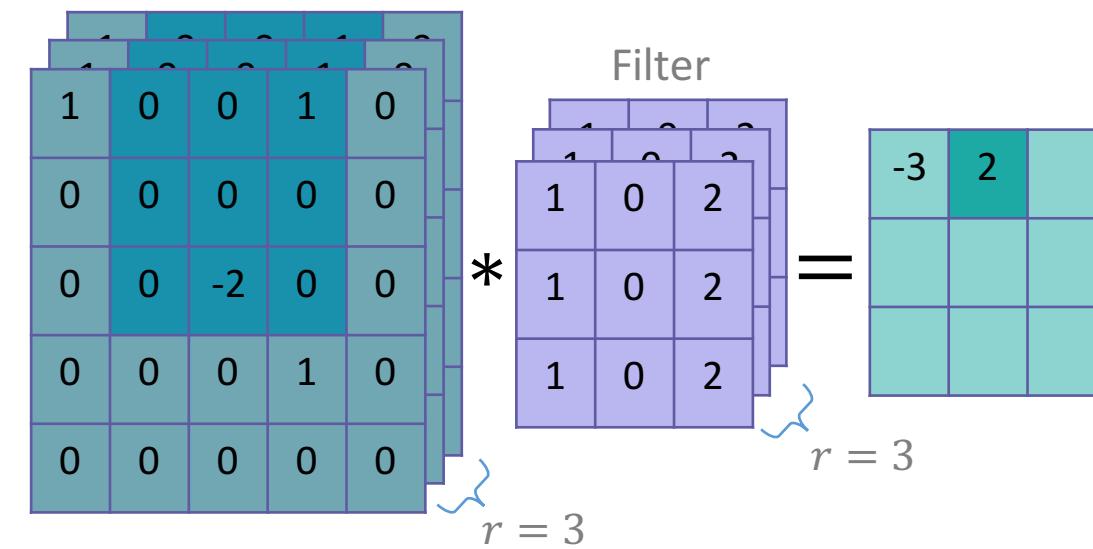
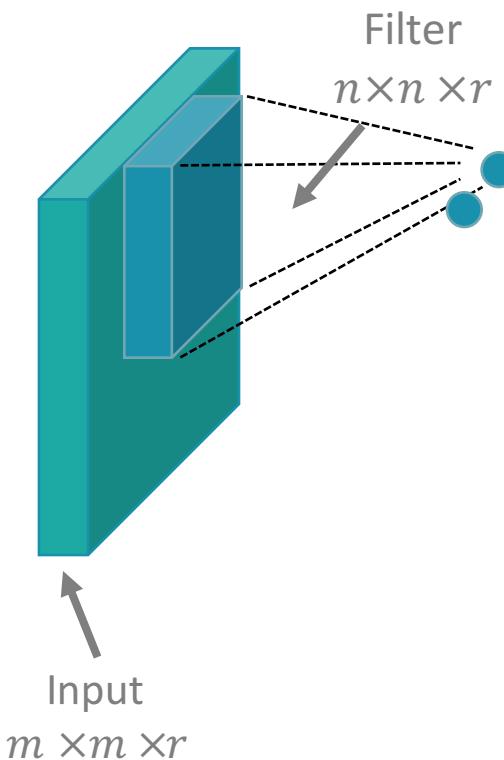
$$(I * F)_{r,s} := \sum_{u=0}^{r-1} \sum_{v=0}^{r-1} I_{r+u,s+v} \cdot F_{u,v}$$



Building Blocks of a CNN: Convolutional Layer

OR HOW TO GET EVEN BETTER CLASSIFICATION RESULTS

Output of the neuron: **discrete convolution** of all input-features for one filter.



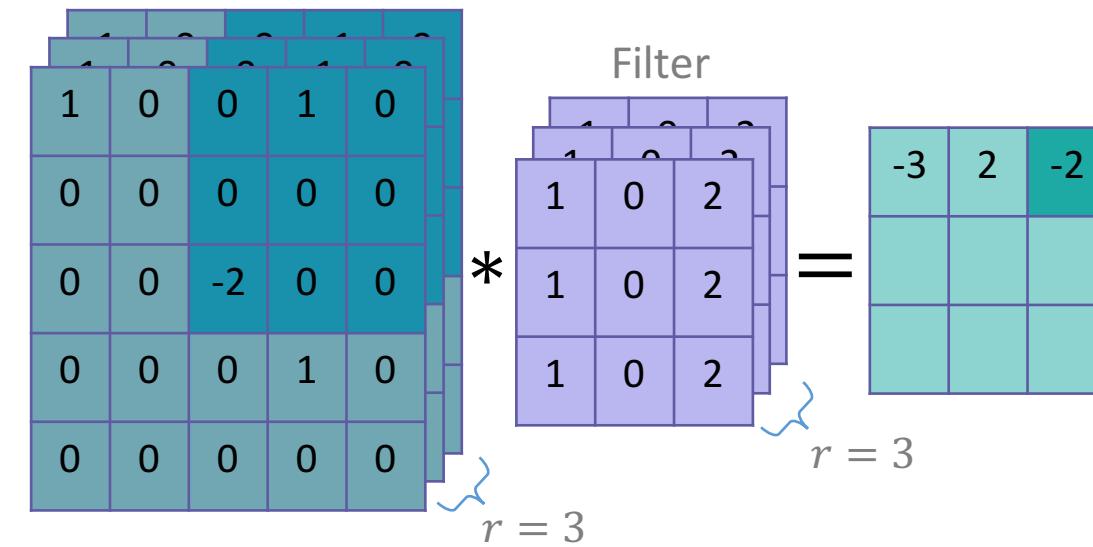
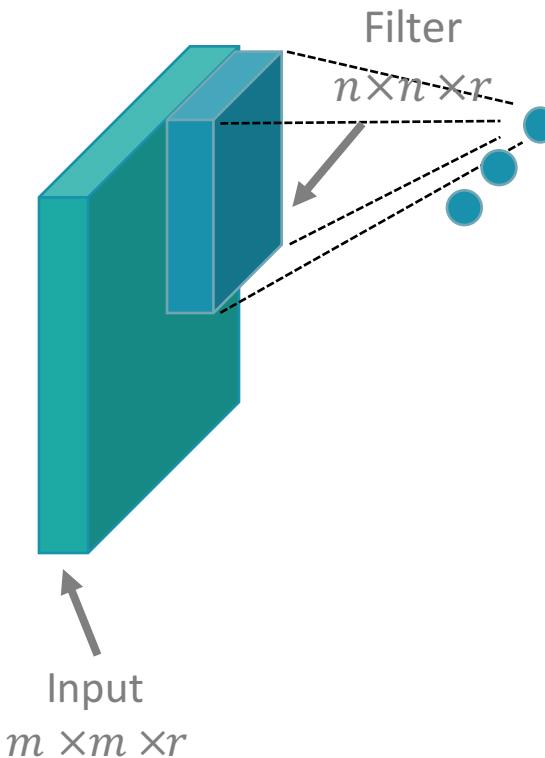
$$(I * F)_{r,s} := \sum_{u=0}^{r-1} \sum_{v=0}^{r-1} I_{r+u,s+v} \cdot F_{u,v}$$



Building Blocks of a CNN: Convolutional Layer

OR HOW TO GET EVEN BETTER CLASSIFICATION RESULTS

Output of the neuron: **discrete convolution** of all input-features for one filter.



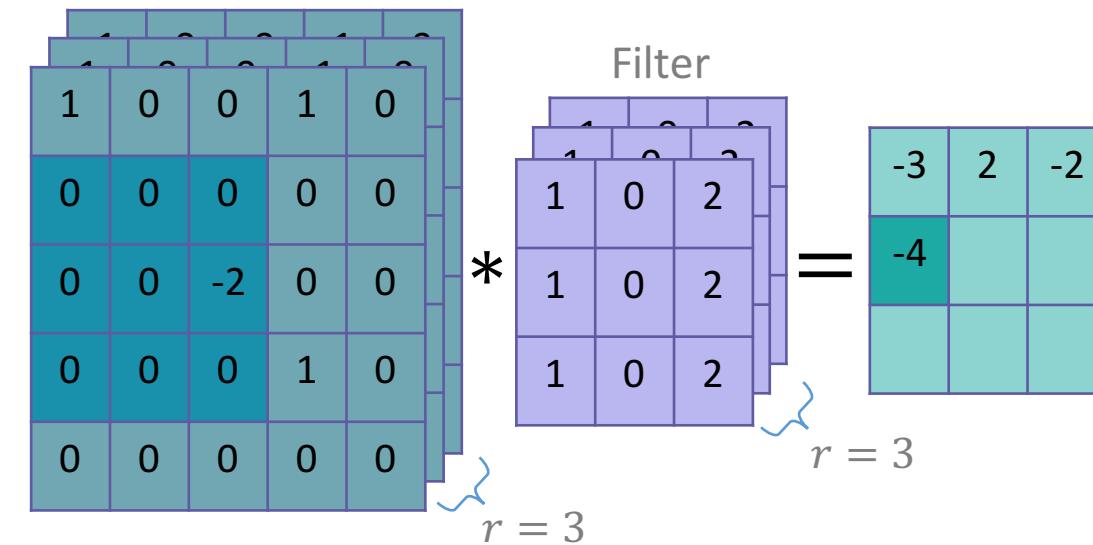
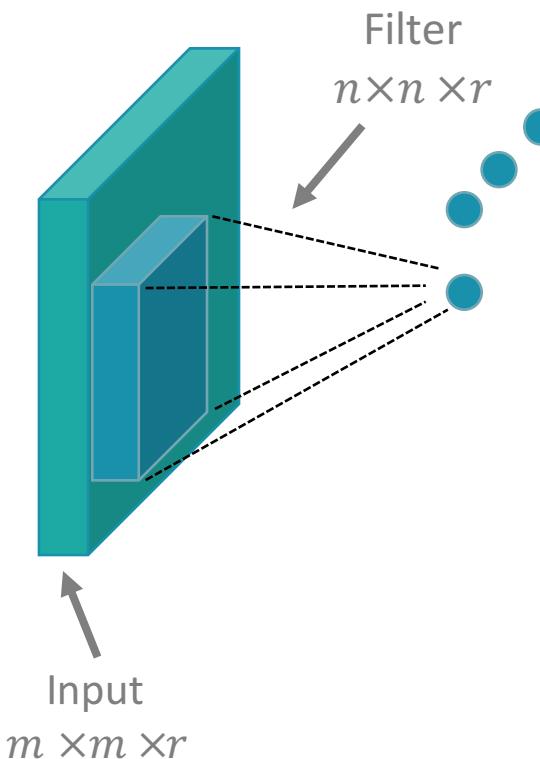
$$(I * F)_{r,s} := \sum_{u=0}^{r-1} \sum_{v=0}^{r-1} I_{r+u,s+v} \cdot F_{u,v}$$



Building Blocks of a CNN: Convolutional Layer

OR HOW TO GET EVEN BETTER CLASSIFICATION RESULTS

Output of the neuron: **discrete convolution** of all input-features for one filter.



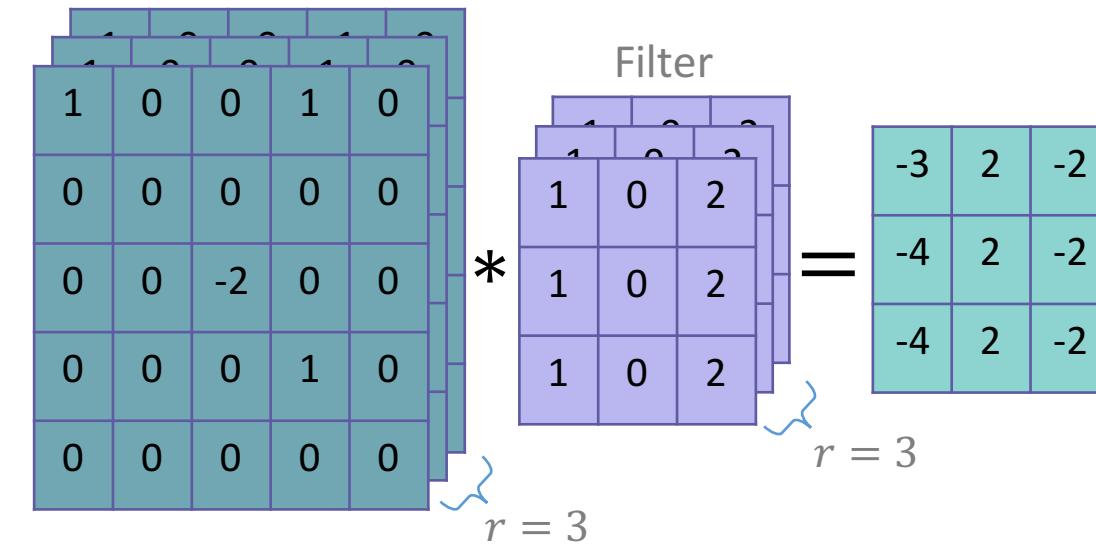
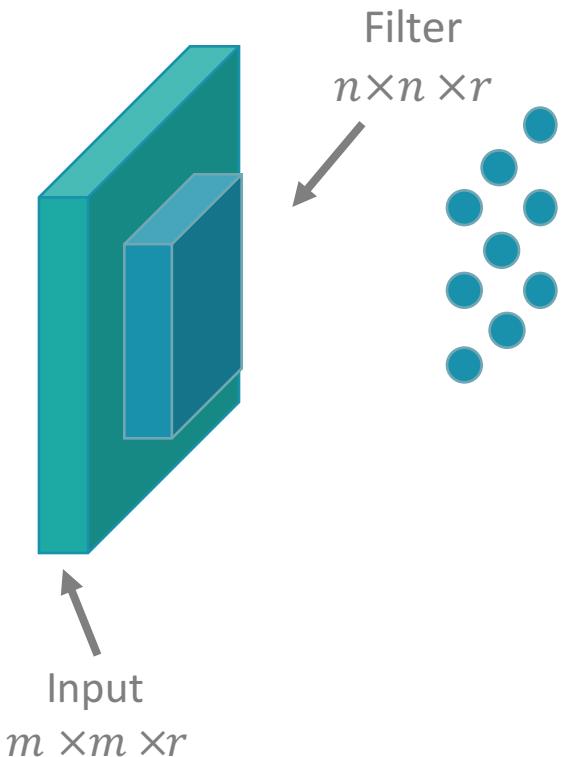
$$(I * F)_{r,s} := \sum_{u=0}^{r-1} \sum_{v=0}^{r-1} I_{r+u, s+v} \cdot F_{u,v}$$



Building Blocks of a CNN: Convolutional Layer

OR HOW TO GET EVEN BETTER CLASSIFICATION RESULTS

Output of the neuron: **discrete convolution** of all input-features for one filter.



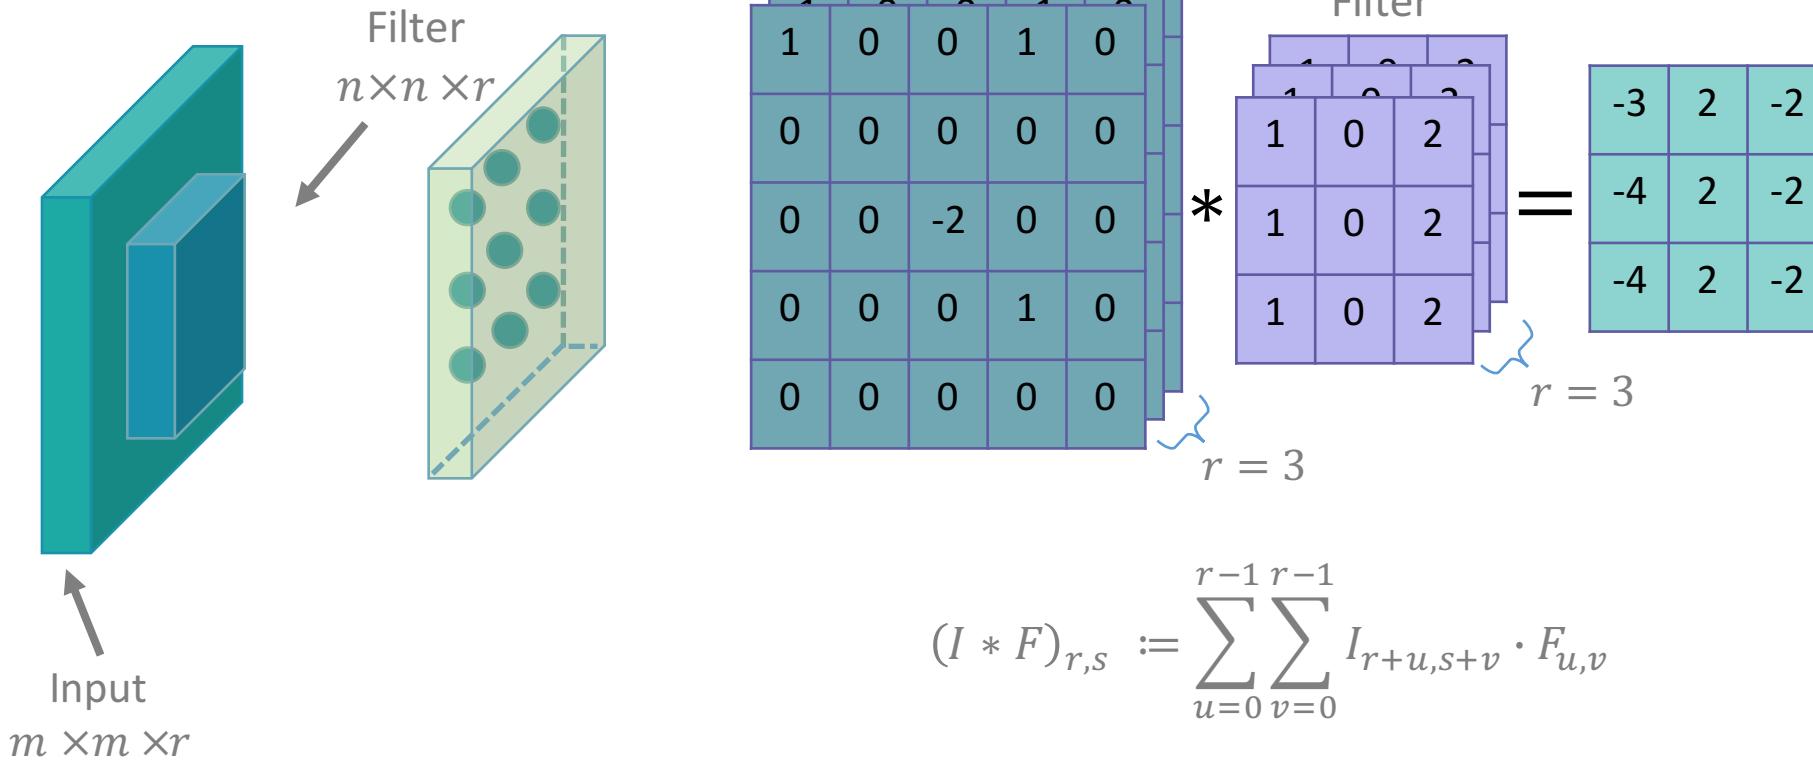
$$(I * F)_{r,s} := \sum_{u=0}^{r-1} \sum_{v=0}^{r-1} I_{r+u,s+v} \cdot F_{u,v}$$



Building Blocks of a CNN: Convolutional Layer

OR HOW TO GET EVEN BETTER CLASSIFICATION RESULTS

Output of the neuron: **discrete convolution** of all input-features for one filter.

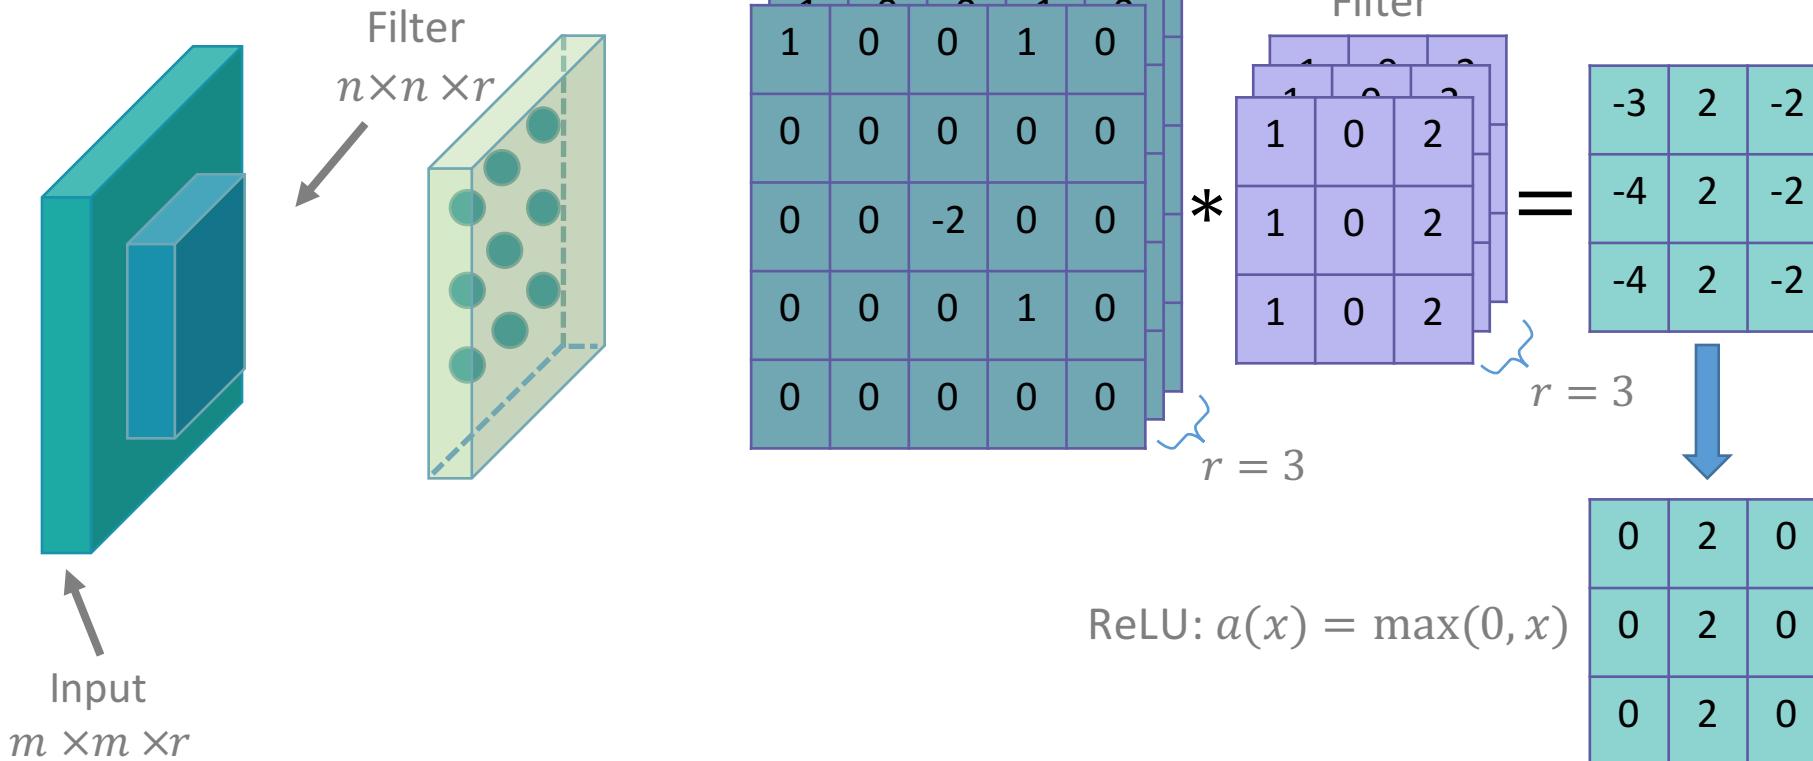


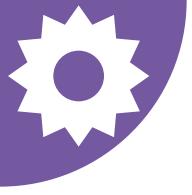


Building Blocks of a CNN: Convolutional Layer

OR HOW TO GET EVEN BETTER CLASSIFICATION RESULTS

Output of the neuron: **discrete convolution** of all input-features for one filter.

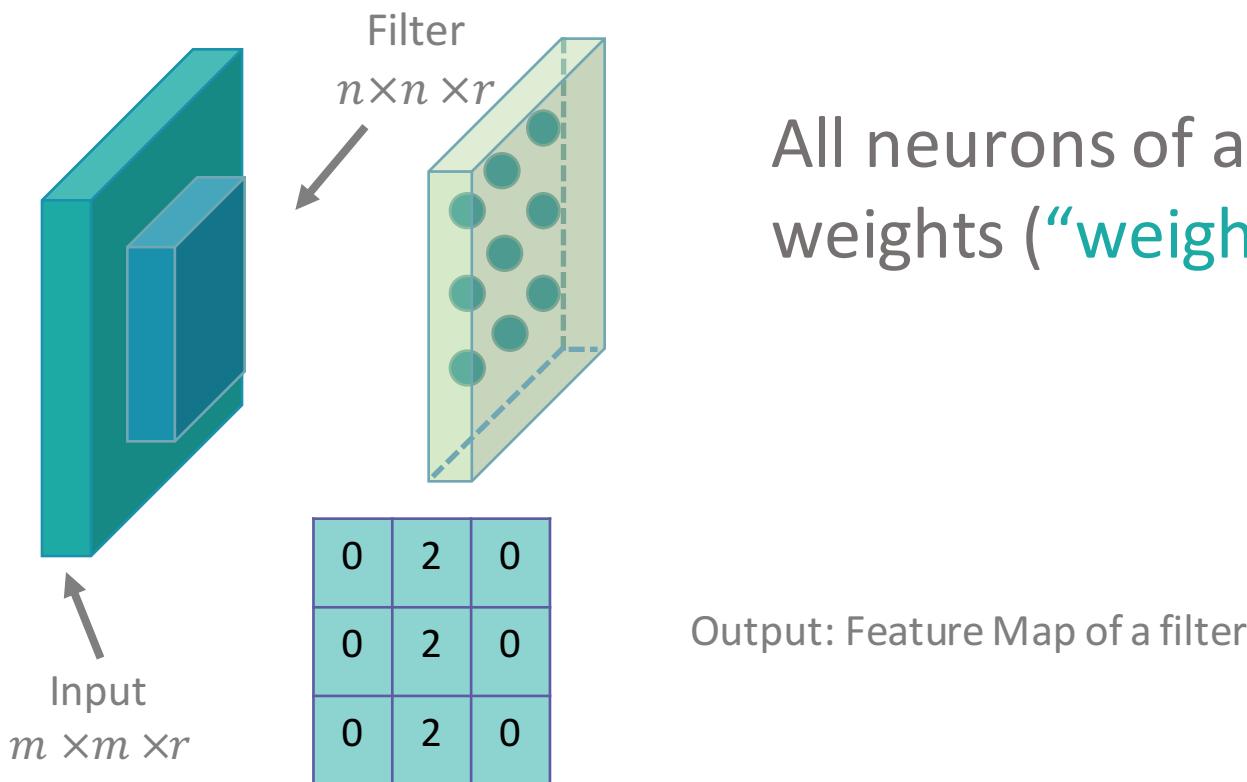




Building Blocks of a CNN: Convolutional Layer

OR HOW TO GET EVEN BETTER CLASSIFICATION RESULTS

Output of the neuron: **discrete convolution** of all input-features for one filter.

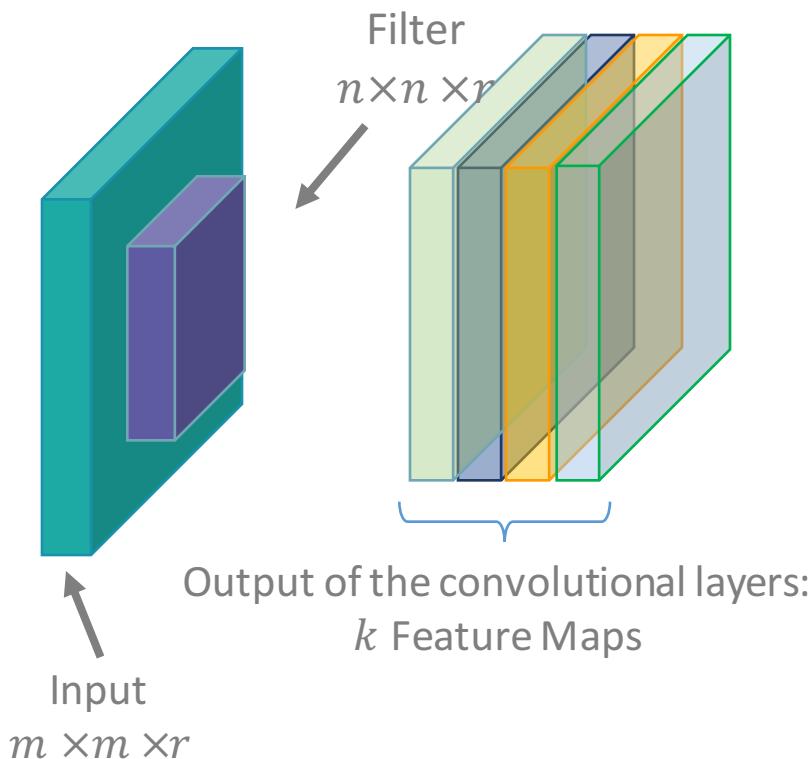


All neurons of a filter share the same weights ("weight sharing")!



Building Blocks of a CNN: Convolutional Layer

OR HOW TO GET EVEN BETTER CLASSIFICATION RESULTS



Input of a convolutional layer:
 r feature maps

Output: k filters create k feature
maps



Simple CNN using TensorFlow on MNIST

OR HOW TO GET EVEN BETTER CLASSIFICATION RESULTS

Live demo
[cnn_mnist.ipynb](#)



Integrated Perception

SEGMENTATION, DETECTION AND CLASSIFICATION – ALL AT ONCE

Application

- Vehicle Detection
- Road Segmentation
- Road Type Classification

Technique

- Fully Convolutional Network
- Supervised Learning

Data

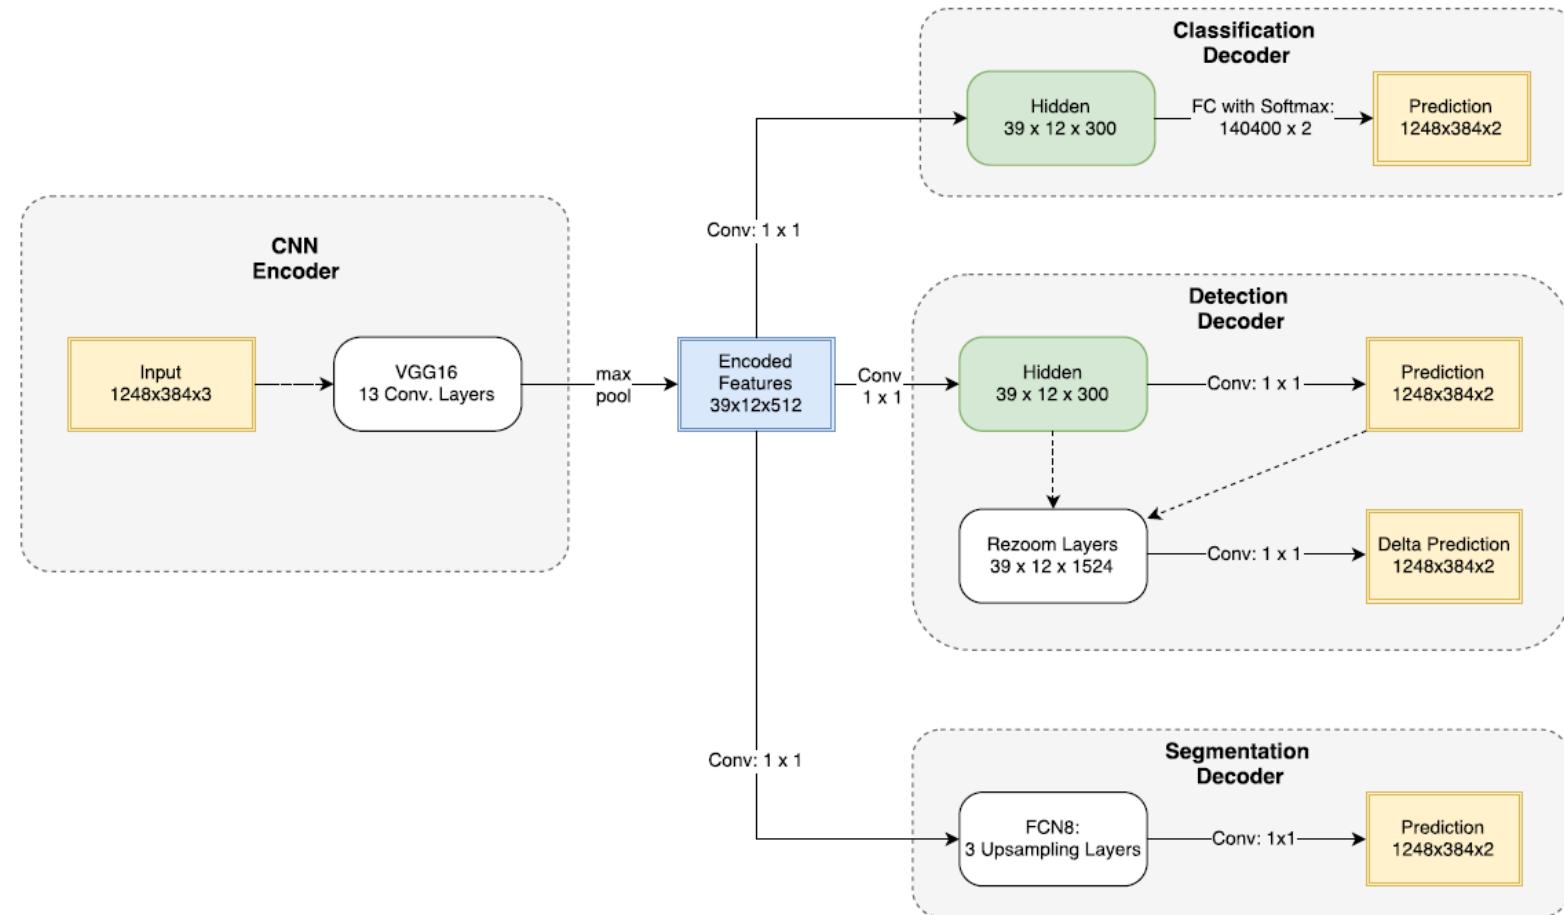
- KITTI Dataset





MultiNet Architecture

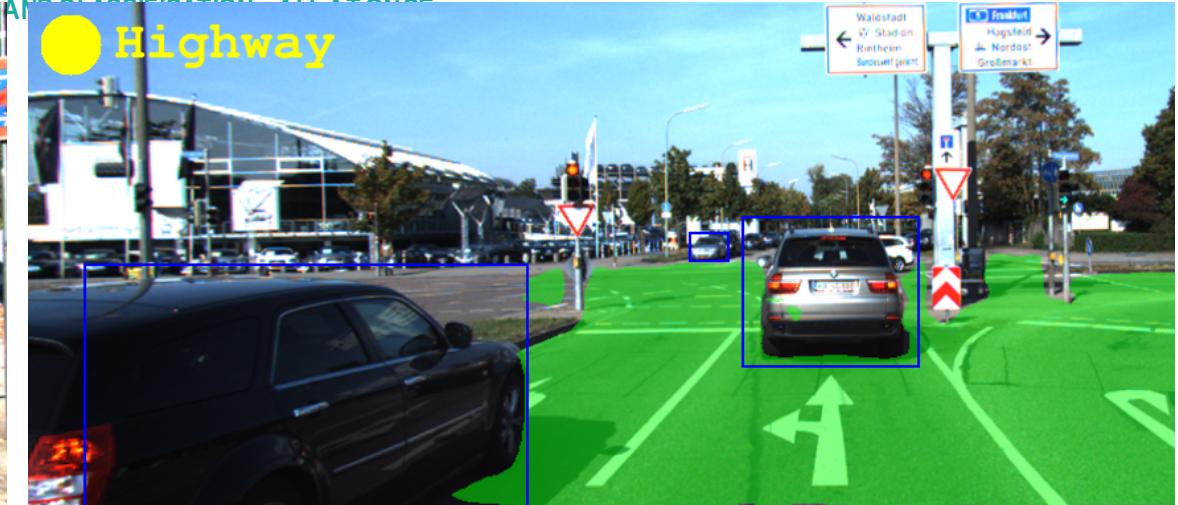
SEGMENTATION, DETECTION AND CLASSIFICATION – ALL AT ONCE





MultiNet Results

SEGMENTATION, DETECTION AND CLASSIFICATION - ALL AT ONCE





DeepTLR

DETECTION AND CLASSIFICATION OF TRAFFIC LIGHTS



CNN to detect traffic lights

- Implicit sliding window approach
- Result contains probability maps and bound box proposals

Post-Processing

- Clustering of bounding boxes to achieve detection

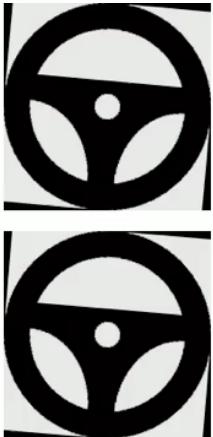
Performance

- Runnable with 33 Hz (640x480 px)
- Precision: 95.6%
- Recall: 91.4%

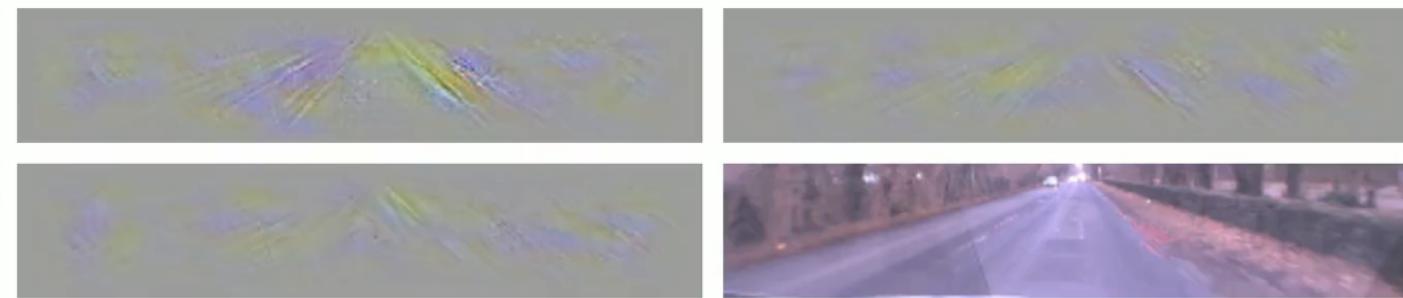


End-to-End Driving

END-TO-END LEARNING FOR AUTONOMOUS DRIVING



Distance lane center: 0.00
Orientation diff: 0.00





TensorFlow Yourself

HOW TO GET INVOLVED?

Online Learning Resources

- TensorFlow Online Tutorial:
https://www.tensorflow.org/get_started/
- Stanford CS231n. Full course available on YouTube:
<https://www.youtube.com/playlist?list=PLkt2uSq6rBVctENoVBg1TpCC7OQi31AIC>
- .. many many more! TensorFlow has a large community and is pretty well documented.

ML-KA Team Tensorflow

- Tuesdays, 5:30 pm
- Dates will be added to Meetup!

ML-KA Einsteiger Track

- Dates will also be added to Meetup!

.. or: Come talk to me!

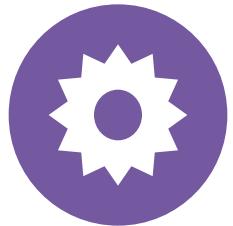
Conclusion



Overview
over
Frameworks



TensorFlow
Basics



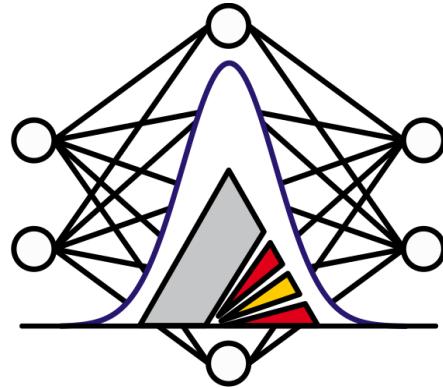
Coding
Examples



TensorFlow
for Autonomous
Driving



How to get
involved



Karlsruhe Machine Learning,
Statistics and AI Meetup

Any questions?